
Challenges for scaling DL models with GPUs
23th October 2023
Sergi Guimerà Roig
sergi.guimera@estudiantat.upc.edu

Cost

Intro

Models

≠

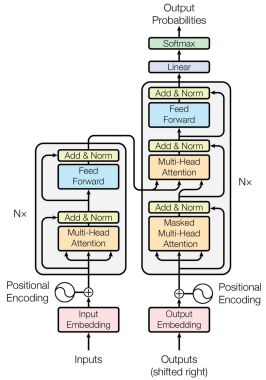

Figure 1: The Transformer - model architecture.

$$\begin{array}{c} \\ 1 \\ 2 \\ 3 \\ \vdots \\ m \end{array} \begin{array}{c} \begin{array}{cccc} 1 & 2 & \cdots & n \end{array} \\ \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ a_{31} & a_{32} & \dots & a_{3n} \\ \vdots & \vdots & \vdots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix} \end{array}$$

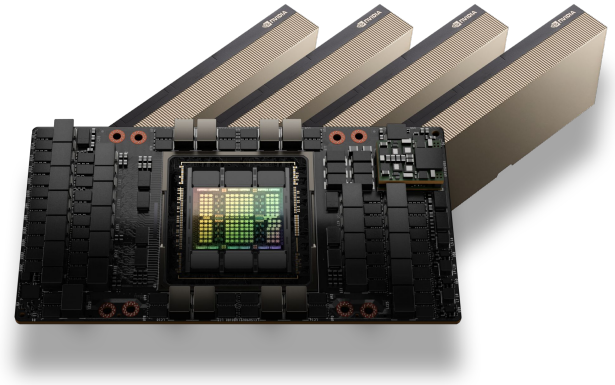

GPU

Scale

AG How many letters are there in the word "nineteen"?

There are eight letters in the word "nineteen".

AG are you sure?

I apologize for my mistake. There are actually nine letters in the word "nineteen". Thank you for pointing that out.

Scalability

A100 ~60% Model/Hardware FLOPS

Costs GPU


FEDERAL RESERVE NOTE
THE UNITED STATES OF AMERICA
THIS NOTE IS LEGAL TENDER FOR ALL DEBTS, PUBLIC AND PRIVATE
B 03542754 F
WASHINGTON, D.C.
ONE
B 03542754 F
ONE DOLLAR


Microsoft Azure
Hello, I like money

State-of-the-art models

Model	Size (# Parameters)	Training Tokens
LaMDA (Thoppilan et al., 2022)	137 Billion	168 Billion
GPT-3 (Brown et al., 2020)	175 Billion	300 Billion
Jurassic (Lieber et al., 2021)	178 Billion	300 Billion
Gopher (Rae et al., 2021)	280 Billion	300 Billion
MT-NLG 530B (Smith et al., 2022)	530 Billion	270 Billion
Chinchilla	70 Billion	1.4 Trillion

State-of-the-art costs

mosaicML

30B → $450k | $326k
70B → $2.5M | $1.75M

PaLM

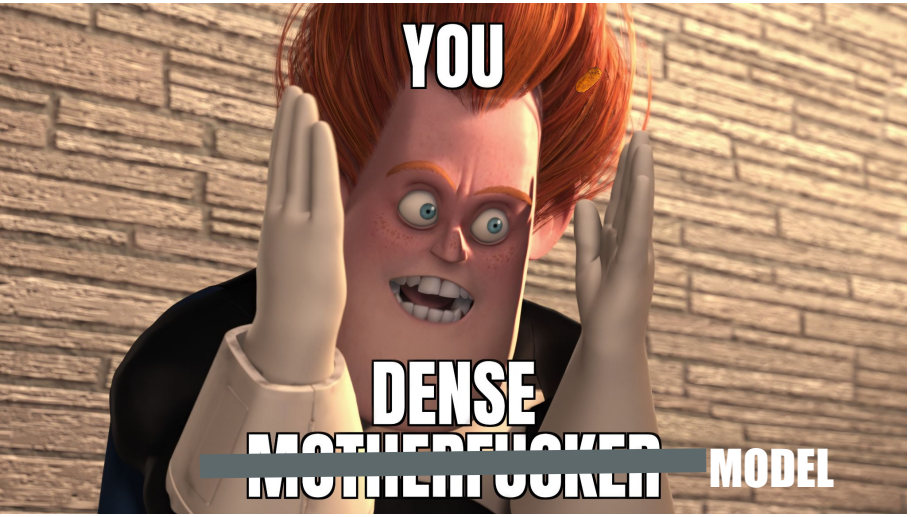

Dense Transformer Scaling Wall

1T parameters → $300M
w/ 100k A100 → 3 months

Not so fast